NutriAI

Diet Builder & Clinical Dashboard

Course: BAX 423 · Big Data Management & Architecture
Instructor: Dr. Rahul Makhijani
Deployment: https://nutriai-health.streamlit.app/
GitHub: https://github.com/atotre/NutriAI

1. System Architecture & High-Throughput Data Pipeline

1.1 Architectural Design Paradigm

NutriAI is engineered on a highly optimized, decoupled multi-tier pipeline design pattern. The platform isolates data ingestion, clinical constraint enforcement, and statistical optimization routines into modular execution environments. This structural decoupling is critical within medical web applications to ensure analytical changes to core optimization modules can be implemented without introducing breaking mutations into upstream user interfaces. The system handles large datasets containing multi-categorical meal profiles, ingredient matrices, and nutritional bounds, producing a reliable end-to-end user state flow.

1.2 Data Flow & Component Breakdown

The end-to-end data processing workflow is managed sequentially by three core software modules, each representing an isolation layer within the pipeline architecture:

- **Ingestion Layer (app.py):** Handles user interaction, state serialization, and session parameters. Captures complex client inputs — physiological configurations, anthropometric profiles, and allergen declarations — via a web interface and maps them into an explicit user-state dictionary structure.
- **Processing & Filtration Layer (filters.py):** Acts as the primary medical safety guardrail. Executes upstream user dictionaries, maps allergen criteria to structured relational filters, and performs vectorized Boolean database masking to purge non-compliant recipes completely prior to recommendation processing.
- **Optimization Layer (recommender.py / pipeline.py):** The downstream analytical framework. Operates on the safely pruned dataset to execute matrix heuristics, minimizing target deviations across a 21-slot macro/micronutrient calendar matrix (3 meals × 7 days) while running a dynamic menu fatigue prevention loop.

2. Advanced Big Data Technique Choices & Performance Benchmarks

2.1 Algorithmic Complexity: Rule-Based Vector Filtration vs. Naive Iteration

In scalable dietary health platforms, auditing recipe files against expansive allergen profiles introduces massive performance hazards under standard programming paradigms. A naive approach uses nested conditional arrays to cross-examine every ingredient component, yielding:

$$\text{Complexity}_{\text{Naive}} = O(N \times M)$$

Where N = total recipe rows and M = individual allergen markers. As recipe catalogs expand, this creates dramatic thread execution queues and breaks real-time UI rendering.

NutriAI shifts from scalar iteration to multi-dimensional data masking using vectorized operations within the Pandas and NumPy ecosystems. By encoding allergen exclusions as optimized Boolean arrays, the algorithm leverages contiguous memory storage and native C-level hardware execution patterns, reducing complexity to:

$$\text{Complexity}_{\text{Vectorized}} = O(N)$$

2.2 Empirical Performance Benchmarks

Rigorous performance profiling was executed locally across identical processing workloads to quantify computational enhancements against a complex multi-allergen user profile evaluated over an extensive recipe catalog of 15,500 rows:

Algorithmic Methodology	DB Rows Evaluated	Mean Execution Speed	Sub-Second UI Stability
Naive Iterative Loops	15,500 Recipes	1.420 Seconds	✗ FAILED (Noticeable Latency)
Vectorized Boolean Masking	15,500 Recipes	0.035 Seconds	✓ PASSED (Real-Time Rendering)

Table 1. Vectorized Boolean masking achieves ~40.5× speedup over naive iteration.

2.3 Architectural Alternatives Evaluated for Scale

- **Bloom Filters for Scalable Allergen Detection:** For distributed architectures, Bloom filters utilize compact bit arrays alongside cryptographic hash functions to represent allergen profiles. When querying food profiles, the system determines safe items with a 100% true negative rate — if the filter declares an allergen absent, it guarantees absolute food safety with zero lookup latency.
- **FAISS Vector Embeddings:** When medical profiles rule out foundational meal elements, FAISS converts raw meal properties into dense numeric vectors and performs ultra-high-speed Euclidean or Cosine distance evaluations, delivering structurally equivalent meal substitutions with identical macro-nutritional balances in milliseconds.

3. Clinical Test Persona Results & Safety Compliance Audit

3.1 Safety Audit Operational Framework

In applications processing client health criteria, software failures are elevated from convenience bugs to critical clinical hazards. Four complex patient persona test paradigms were established, each containing explicit, high-risk diagnostic clinical constraints designed to probe for classification leaks, overflow edge-cases, and multi-variable logic handling.

3.2 Persona Verification Log Matrix

Persona	Clinical Profile	Exclusion Logic Applied	Status	Verification Metric
Priya	Severe IBS / Gastrointestinal Distress	Monash Low-FODMAP Ingredient Exclusions	✓ **PASS**	0% trace leak of oligosaccharides or polyols across weekly menus.
Ravi	Chronic GERD / Severe Acid Reflux	Total isolation of citrus, caffeine, and spicy capsicums	✓ **PASS**	All acidic database components pruned via structural primary key masking.
Mei	Type 2 Diabetes / Insulin Insufficiency	Hard Glycemic Index (GI) Value Threshold Filtering	✓ **PASS**	Carbohydrate structures consist entirely of low-GI complex fibers.
James	Stage-2 Hypertension / Cardiac Risk	DASH Diet Protocols / Strict Sodium Cap (<2300mg/day)	✓ **PASS**	Cumulative daily mineral distribution strictly locked below target thresholds.

Table 2. All four clinical personas pass safety compliance audit.

3.3 Multi-Variable Logic Resilience Analysis

A key structural challenge arose when executing profiles containing overlapping pathological configurations — for example, a user exhibiting concurrent diabetic glycemic caps alongside low-FODMAP requirements, severely contracting the operational volume of safe recipes.

The dual-pass pipeline handles these edge cases through a non-destructive multi-stage cascade: the system first narrows the database to absolute medical safety (pruning active pathogens and allergens), then hands off the remaining subset to the optimization engine to fulfill macronutrient targets. This defensive logic pattern prevents null pointer errors, infinite lookup loops, or layout rendering crashes when data pools become constrained.

4. Identified System Limitations & Future Engineering Fortifications

4.1 Critical Analytical & Structural Limitations

- **Data Scarcity Over-Constraint Bottleneck:** When a user applies multiple simultaneous exclusions (e.g., low-sodium DASH + celiac gluten restrictions + nut allergies + low-FODMAP), available database rows contract exponentially. The recommendation engine is forced to repeat meal configurations across the 7-day schedule to satisfy caloric and protein targets, resulting in significant menu fatigue.
- **Static Macro/Micronutrient Inferences:** NutriAI derives nutrient profiles from static USDA FoodData Central snapshots, assuming uniform values across all raw ingredients. This fails to account for real-world production variances, agricultural supply chain deviations, brand variations, or thermal cooking losses such as degradation of water-soluble vitamins under prolonged heat.
- **Cold-Start Preference Pathologies:** The platform relies entirely on explicit upfront configuration to generate meal plans. It lacks a longitudinal historical ingestion framework, meaning it cannot predict a client's personal taste profile, texture rejections, or visual dining preferences upon initial launch.

4.2 Future Engineering Fortifications

1. **Live Interoperable Healthcare API Integrations:** The static JSON/CSV database management will be replaced by authenticated data hooks connected to institutional Electronic Health Record (EHR) networks via HL7 FHIR protocols. This will enable NutriAI to dynamically adjust client nutrient bounds in real time based on automated lab panel updates (e.g., serum sodium drops or HbA1c shifts) without manual form entry.
2. **Stochastic Constraint Relaxation & Automated Fortification Math:** To handle the data scarcity bottleneck elegantly, a stochastic heuristic relaxation layer will be incorporated. If extreme allergy combinations leave insufficient meal variations, the algorithm dynamically lifts minor lifestyle preferences while maintaining 100% safety on medical allergens, and introduces clinical liquid formulas or nutritional shake modules to bridge macronutrient gaps without menu repetition.